


RETAIL BROKERAGE · COMPETITIVE INTELLIGENCE

Competitive Intelligence Dashboard

Scrapes competitor fees, scores public sentiment, and ranks every broker on a single vulnerability index — so the question stops being “what do they charge?” and becomes “which one do we go after?”

Python · Streamlit · FinBERT · pdfplumber · pandas



Every competitor lands in one of four quadrants.

THE PROBLEM

Price tells you half the story

1

Fees are buried in awkward places

Public pricing pages for the neobrokers, 40-page German Preis- und Leistungsverzeichnis PDFs for the incumbents. Nobody re-reads those monthly.

2

Sentiment lives somewhere else entirely

Reddit threads, news headlines, app-store complaints. Genuinely useful, but not comparable to a euro figure sitting in a table.

3

Neither one, alone, triggers a decision

A fee table and a mood summary still leave you guessing which competitor is actually exposed right now.

THE APPROACH

One number per competitor

Standardise both signals against the peer set, weight them, and rank. A high score means a broker charges more than its peers and is liked less than its peers at the same time.

Competitive Vulnerability Score > 1.0 = go take their customers

Two pipelines, one engine

Fee ingestion

`Unified_Ingestion.py`

CSS selectors on pricing pages · pdfplumber
table scan with German fee keywords

Sentiment

`sentimental_analysis.py`

FinBERT over Reddit (r/Finanzen, r/investing)
and Google News RSS → -1.0 to 1.0



Strategy engine

`strategy_n_CVS.py`

- Z-score both signals
- Weight and combine → CVS
- Map to a tactic



Streamlit dashboard

`Dashboard.py`

Plotly quadrant chart, live weight
sliders, ranked tactics table

Current state: the two pipelines run standalone and the dashboard reads sample data. Connecting them end-to-end is the next commit.

THE MODEL


Standardise, weight, rank

```
z_price      = (fee - mean) / std
z_sentiment  = (score - mean) / std
CVS = wp · z_price - ws · z_sentiment
```

Euros and a -1 to 1 mood scale can't be added directly, so both are converted to distance-from-peer-average first. Weights default to 0.6 / 0.4 and are live sliders — dial price up when a rate war starts.

- **CVS > 1.0** Aggressive acquisition — cost-leadership marketing
- **CVS < -1.0** Avoid head-on — differentiate on niche features
- **In between** Quadrant rules: rehabilitate, defend margins, or monitor

Competitive Vulnerability Score



Sample peer set shipped with the repo. Scores are relative — add a broker and everyone moves.

Where it stands

Working today

- ✓ Z-score engine and tactic rules — deterministic, no I/O, testable
- ✓ FinBERT scoring across Reddit and Google News headlines
- ✓ PDF fee extraction with German keyword mapping
- ✓ Streamlit UI with live weight sliders and a quadrant chart

Not wired up yet

- Dashboard reads sample data — the pipelines aren't plugged in
- Scraper selectors and the PLV PDF URL are placeholders
- X/Twitter source returns mock posts; needs paid API access
- On a 3-broker set the ± 1.0 thresholds rarely fire — the obvious target scored 0.81 and still read “monitor”

Next: run ingestion → sentiment → engine end-to-end against a real 8-broker peer set, which is also what the thresholds need to behave.